

# Sistemas Distribuídos

José Orlando Pereira

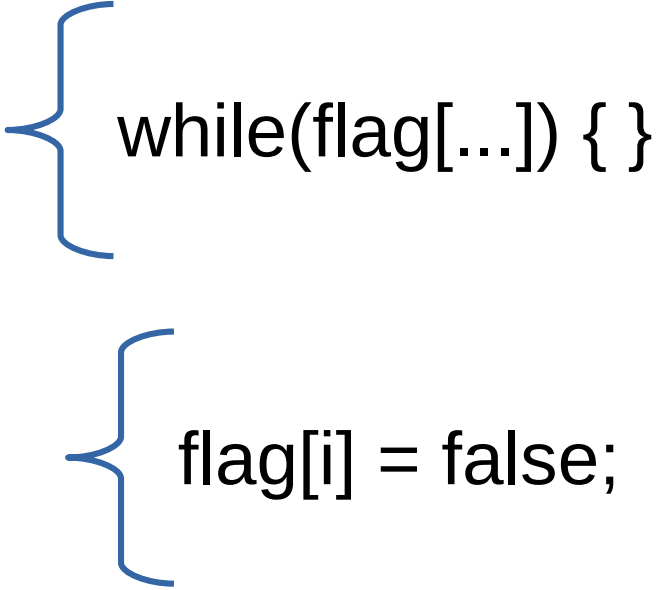
Departamento de Informática  
Universidade do Minho



# Mutex with Peterson's/...

- Does it really work?

```
try {  
    l.lock();  
    c=c+1;  
} finally {  
    l.unlock();  
}
```





while(flag[...]) { }

flag[i] = false;

# Consequence

```
try {  
    l.lock();  
    c=c+1;  
} finally {  
    l.unlock();  
}
```

 while(flag[...]) { }


 flag[i] = false;

- Initially c=10
- One thread:
  - read c
  - write c = 11
  - write flag[i] = false
- Other thread:
  - read flag[...] = false
  - read c = 10!!!!



# Solution: Memory barriers

- Declare: volatile int j; or AtomicInteger j;
- Reading from a volatile j waits for all writes preceding the observed value on j to be also visible
  - Writer code:
    - i=1; j=1;
  - Reader code:
    - rj=j; ri=i; System.out.println(rj+", "+ri);



waits for write i=1  
to be visible

# Should I use it?

- Synchronization operations also act as memory barriers:
  - `java.util.concurrent.*` primitives
  - `synchronized` keyword
- Volatile variables also impact performance:
  - A volatile access takes as much time as using a lock
  - Therefore... **use the lock!**

# The real world

- Is j.u.c.ReentrantLock implemented with Bakery or something similar?
  - **No!**
  - Needs  $O(n)$  space for each lock (with  $n$  threads)
  - Consumes CPU time / power when busy
- These algorithms are a good for:
  - Understanding concurrent programming and races
  - Preparing for distributed algorithms

# The real world

- `j.u.c.ReentrantLock` is implemented with atomic operations and scheduling:
  - Atomic operations (`testAndSet`, `compareAndSwap`, ...) are used to quickly check if mutex is available
    - For more on atomic operations, see Chapter 7 of TAOMP (not part of the program!)
  - The thread is suspended by the OS when waiting for longer periods
    - Remember scheduling from Operating Systems
    - `man futex` on Linux

# j.u.c Locks vs Monitors

```
class C {  
    private int i;
```

There is a hidden “lock” in each object used by “synchronized”

```
synchronized public void m() {
```

```
    i++;
```

```
}
```

```
}
```

```
class C {  
    private int i;  
    private Lock l =  
        new ReentrantLock();
```

```
public void m() {  
    try { l.lock();  
        i++;  
    } finally { l.unlock(); }
```

```
}
```

```
}
```

Equivalent code  
(approximately...)



# j.u.c. Locks vs Monitors

- Main differences, for now:
  - Synchronized blocks are nested in LIFO order  
vs.  
j.u.c. Locks can be unlocked in any order
    - To take advantage of two phase locking
  - Threads waiting for a synchronized block enter in any order  
vs.  
j.u.c. Lock can be configured for threads waiting to enter in a fair (not FIFO!) order

# j.u.c Conditions vs Monitors

- class C {

There is a hidden “condition” in each object used by “wait()/notifyAll()”

```
synchronized public void m1() {  
    while(...) wait();  
}
```

```
synchronized public void m2() {  
    notifyAll();  
}
```

- class C {

```
private Lock l =  
    new ReentrantLock();  
private Condition c =  
    l.newCondition();  
public void m1() {  
    try { l.lock();  
        while(...) c.await();  
    } finally { l.unlock(); }  
}  
public void m2() {  
    try { l.lock();  
        c.signalAll();  
    } finally { l.unlock(); }  
}
```

Equivalent code  
(approximately...)

# j.u.c. Conditions vs Monitors

- Main differences:
  - One implicit condition for each lock  
vs.  
Many j.u.c. conditions for the same lock
    - Avoids signalAll()
  - Threads waiting for a condition wakeup in any order  
vs.  
Threads waiting for a j.u.c. Condition obtained from a ReentrantLock wakeup in FIFO order (but may not acquire lock in FIFO order...)